

AMS 691.03: Deep Learning (Fall 2023)

Programming Project #2 Due December 8th, 2023

1. You need to submit (1) a report in PDF and (2) your code files, both to Brightspace.
2. Your PDF report should include (1) answers to the non-programming part, and (2) results and analysis of the programming part. For the programming part, your PDF report should at least include the results you obtained, for example the accuracy, training curves, parameters, etc. You should also analyze your results as needed.
3. Please put all your files (PDF report and code files) into a compressed file named “PP#_FirstName_LastName.zip”
4. Unlimited number of submissions are allowed on Brightspace and the latest one will be timed and graded.
5. Please read and follow submission instructions. No exception will be made to accommodate incorrectly submitted files/reports.
6. All students are highly encouraged to typeset their reports using Word or L^AT_EX. In case you decide to hand-write, please make sure your answers are clearly readable in scanned PDF.
7. Only write your code between the following lines. Do not modify other parts.
YOUR CODE HERE
END YOUR CODE

Recurrent Neural Network (RNN) for Language Modeling: In this assignment, you will implement a recurrent neural network (RNN) for language modeling using Pytorch. The task is to predict word x_{t+1} given words x_1, \dots, x_t :

$$P(x_{t+1} = v_j | x_t, \dots, x_1)$$

where v_j is the j -th word in the vocabulary. The file “utils.py” gives an example of how to generate the vocabulary. You can read it if interested. With the vocabulary, we can transform a word x_i into a one-hot vector.

Our RNN model is, for $t = 1, \dots, n - 1$:

$$\begin{aligned} e^{(t)} &= x^{(t)} L, \\ h^{(t)} &= \text{sigmoid}(h^{(t-1)} H + e^{(t)} I + b_1), \\ \hat{y}^{(t)} &= \text{softmax}(h^{(t)} U + b_2), \\ \bar{P}(x_{t+1} &= v_j | x_t, \dots, x_1) = \hat{y}_j^{(t)}. \end{aligned}$$

where the first line actually corresponds to a word embedding lookup operation. $h^{(0)}$ is the initial hidden state, $\hat{y}^{(t)} \in \mathbb{R}^{|V|}$ and its j -th entry is $\hat{y}_j^{(t)}$.

Training parameters θ in this model are:

- L – embedding matrix which transforms words in the vocabulary into lower dimensional word embedding.
 - H – hidden transformation matrix.
 - I – input transformation matrix which takes word embedding as input.
 - U – output transformation matrix which projects hidden state into prediction vector.
 - b_1 – bias for recurrent layer.
 - b_2 – bias for projection layer.
1. Let the dimension of word embedding as d , the size of vocabulary as $|V|$, the number of hidden units as D , please provide the size of each training parameter above.
 2. To evaluate a language model, we use *perplexity*, which is defined as the inverse probability of the target word according to the model prediction \bar{P} :

$$PP^{(t)}(y^{(T)}, \hat{y}^{(t)}) = \frac{1}{\bar{P}(x_{t+1}^{pred} = x_{t+1} | x_t, \dots, x_1)} = \frac{1}{\sum_{j=1}^{|V|} y_j^{(t)} \hat{y}_j^{(t)}}.$$

Show the relationship between *cross-entropy* and *perplexity*.

3. Read the starting code very carefully and implement the above model by completing `RNNLM.py`. You only need to code some parts related to the RNN model, such as the architecture, loss function. Follow the instructions in the starting code to understand which parts need to be filled in. When you are done, run `python RNNLM.py`. The starting code supports GPU computation and also provides guidelines on how to support CPU. It is very straightforward but note that CPU computation is slow! You should change the hyperparameters to explore the best configurations. Submit the code with your best hyperparameters. Also report the best testing perplexity score. This model is a *generative* model. At the end of the run, it uses the trained language model to generate sentences with the start words you provide. Be creative and report any results that you think are interesting. Please summarize your results and observations.